

Embedding Space Exploration

PRASANNA KOIRALA · AIAC 536 NATURAL LANGUAGE PROCESSING · KATHMANDU UNIVERSITY

Introduction

The previous note explained *what* word embeddings are. This one **explores them** — taking real vectors and measuring what is actually inside the space.

All numbers below are measured from **GloVe** (glove-wiki-gigaword-100): 400,000 English words, each represented by 100 numbers.

1 Cosine Similarity Between Words

Cosine similarity measures the *angle* between two vectors:

$$\cos(A, B) = (A \cdot B) / (|A| \times |B|)$$

Worked by hand

Real vectors have 100 numbers, which is impossible to do on paper. With 2 numbers each the arithmetic is easy to follow, and it is **exactly the same calculation**.

$$\text{Let } A = [2, 3] \quad B = [3, 1]$$

$$A \cdot B = (2 \times 3) + (3 \times 1) = 6 + 3 = 9$$

$$|A| = \sqrt{(2^2 + 3^2)} = \sqrt{13} = 3.61$$

$$|B| = \sqrt{(3^2 + 1^2)} = \sqrt{10} = 3.16$$

$$\cos(A, B) = 9 / (3.61 \times 3.16) = \mathbf{0.79} \quad \rightarrow \text{similar}$$

$$\text{Now } C = [-3, 1]$$

$$A \cdot C = (2 \times -3) + (3 \times 1) = -6 + 3 = -3$$

$$|C| = \sqrt{((-3)^2 + 1^2)} = \sqrt{10} = 3.16$$

$$\cos(A, C) = -3 / (3.61 \times 3.16) = \mathbf{-0.26} \quad \rightarrow \text{unrelated}$$

Measured on real GloVe vectors

Four groups of words were chosen. If the embedding space is meaningful, words *within* a group should score high and words *across* groups should score low.

WITHIN A GROUP	COS	ACROSS GROUPS	COS
cat — dog	0.88	cat — computer	0.30
computer — software	0.84	king — internet	0.25
king — prince	0.77	nepal — horse	0.09
king — queen	0.75	queen — software	0.04

The pattern holds cleanly: **within-group scores are 0.75–0.88, across-group scores are 0.04–0.30**. The space genuinely encodes meaning.

BUT LOOK CLOSER — ONE RESULT BREAKS THE PATTERN

nepal — india **0.68** both South Asian neighbours

nepal — france **0.22** both countries, yet *low*

Both pairs are “two countries”, so a category-based system would score them the same. The embedding does not, because it learned from **how the words are used**, not from a category list. Nepal and India appear in the same news stories; Nepal and France rarely do.

Embeddings capture usage, not taxonomy.

2 Vector Arithmetic

Relationships between words appear as **directions**, so they can be added and subtracted. To answer “man is to king as woman is to *what?*”:

```
vector(king) - vector(man) + vector(woman) = ?
```

then find the word whose vector is nearest to the result.

ARITHMETIC	NEAREST WORD	SCORE	RELATIONSHIP
king - man + woman	queen	0.77	gender
paris - france + nepal	kathmandu	0.81	country → capital
walking - walk + swim	swimming	0.80	verb tense
mice - mouse + cat	cats	0.77	singular → plural

Four different kinds of relationship — gender, geography, grammatical tense and plurals — all recovered by the same simple arithmetic. Nobody programmed any of them in.

WHERE IT FAILS

ARITHMETIC	RESULTS
bigger - big + small	larger 0.89 · smaller 0.87 · large 0.79

The expected answer is *smaller*. It is found — but only in **second place**. The top answer, *larger*, keeps the comparative form but loses the reversal from big to small.

This is the same weakness seen in the previous note: the embedding places **antonyms very close together**, because opposites appear in the same contexts. Analogies work often, but not reliably.

A MORE SERIOUS PROBLEM — BIAS

doctor - man + woman **nurse 0.77** · physician 0.72 · doctors 0.68

The correct answer is *doctor* — the profession does not change with gender. The model returns **nurse**.

This is not a bug in the arithmetic. The vectors were trained on human-written text, and that text associates women with nursing and men with doctoring. **The embedding faithfully learned a human bias.** Any system built on these vectors — CV screening, search, translation — inherits it.

3 Visualising the Space with PCA

Each word has 100 numbers, so the space cannot be drawn. **PCA** (Principal Component Analysis) squashes 100 dimensions down to 2 by keeping the directions along which the words differ most.

30 words were chosen from five clear categories — animals, royalty, countries, technology and food — and projected to 2 dimensions.




IMPORTANT CAVEAT

These 2 dimensions capture only **35% of the total variance**. The other **65% of the information cannot be drawn**. Two words that look close on the page may be far apart in the real 100-dimensional space. The picture is a useful summary, never proof.

4 Visualising with t-SNE

t-SNE solves the same problem differently. PCA is **linear** — it finds the flattest sensible shadow of the data. t-SNE is **non-linear** and optimises for one thing only: keeping each word close to its true nearest neighbours.



	PCA	T-SNE
Type	linear	non-linear
Preserves	overall variance	near neighbours
Distance between clusters	meaningful	not meaningful
Same result every run	yes	no — depends on settings
Clusters look	looser	tighter, better separated

t-SNE produces visibly cleaner groups. But its *gaps* must not be read as distances — a wide gap between two clusters does not mean they are especially unrelated.

5 Interpreting the Clusters

Both plots show the same five groups separating on their own. **Nobody labelled these words**. The grouping emerged purely from how the words are used in text — strong evidence that the embedding space encodes meaning.

Two finer patterns are visible inside the clusters:

- **Countries sub-group by region.** Nepal and India sit together, away from France and Brazil — matching the 0.68 versus 0.22 scores from §1.
- **Royalty is extremely tight.** king, queen, prince, princess, throne and crown almost overlap, because they appear in nearly identical contexts.

THE ANOMALY — ONE WORD IN THE WRONG PLACE

“apple” does not sit with the food cluster. In both plots it drifts towards *computer*, *software* and *laptop*.

The reason is that *apple* has **two distinct meanings** — the fruit and the technology company — and GloVe gives every word **exactly one vector**. That single vector has to serve both senses, so it lands in a compromise position between the two clusters, belonging properly to neither.

This is the fundamental limitation of **static** embeddings: *one word, one vector, regardless of context*. It is precisely the problem that contextual embeddings such as BERT were built to solve — they give *apple* a different vector in “apple pie” than in “Apple released a phone”.

Conclusion

The embedding space is genuinely meaningful: similarity scores separate related from unrelated words, arithmetic recovers gender, geography, tense and plural relationships, and five categories cluster without supervision.

But the exploration also exposed three real limits. Analogies are **unreliable** (*bigger* – *big* + *small* ranks *larger* above *smaller*). The vectors carry **human bias** (*doctor* – *man* + *woman* returns *nurse*). And a single vector per word **cannot represent two meanings**, which is why *apple* sits stranded between food and technology.